

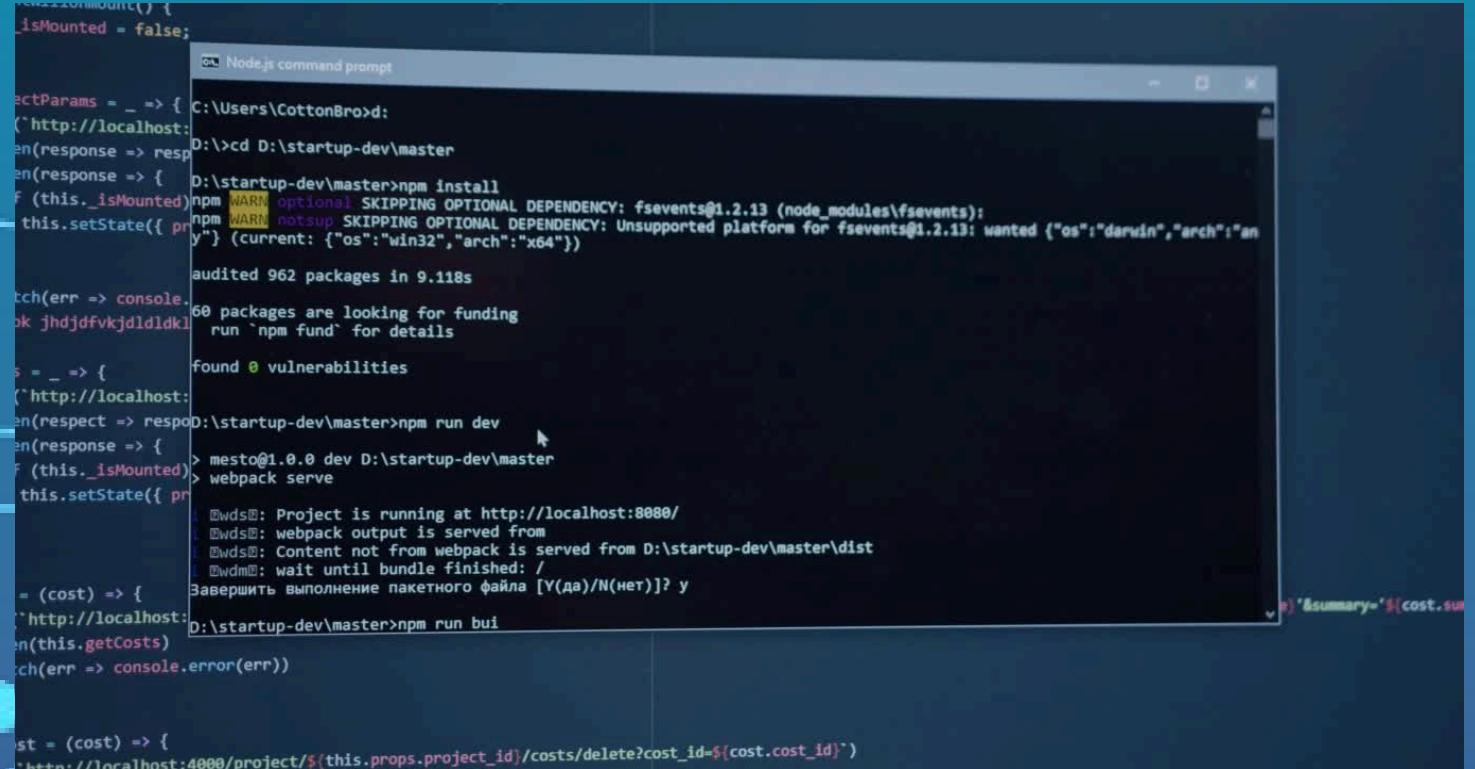
NETWORK INTRUSION DETECTION

Piotr Bonar, Maddox Hurlbert, Sonia Serra
Universitat Autònoma de Barcelona

(group 20)

MOTIVATION

When networks fail... Why detection matters?



```
const component = () => {
  isMounted = false;

  selectParams = _ => {
    C:\Users\CottonBroad: ~
    ^ http://localhost:3001
    en(response => response)
    en(response => {
      D:\startup-dev\master>npm install
      npm WARN optional SKIPPING OPTIONAL DEPENDENCY: fsevents@1.2.13 (node_modules\fsevents):
      npm WARN notsup SKIPPING OPTIONAL DEPENDENCY: Unsupported platform for fsevents@1.2.13: wanted {"os":"darwin","arch":"any"} (current: {"os":"win32","arch":"x64"})
      audited 962 packages in 9.118s
      60 packages are looking for funding
      run `npm fund` for details
      found 0 vulnerabilities
    })
    en(respect => response)
    en(response => {
      if (this._isMounted) {
        mesto@1.0.0 dev D:\startup-dev\master
        > webpack serve
        D:\dms: Project is running at http://localhost:8080/
        D:\dms: webpack output is served from D:\startup-dev\master\dist
        D:\dms: Content not from webpack is served from D:\startup-dev\master\dist
        D:\dms: wait until bundle finished: /
      }
      this.setState({ project: response })
    })
    en(cost => {
      const { project } = this.state
      const { project_id } = project
      const { costs } = project
      const cost_id = costs[cost.id]
      const cost = costs[cost_id]
      const cost_index = costs.indexOf(cost)
      const newCosts = [...costs]
      newCosts[cost_index] = cost
      newCosts[cost_index].isDeleted = true
      newCosts[cost_index].deletedAt = Date.now()
      newCosts[cost_index].cost_id = cost_id
      const newProject = { ...project, costs: newCosts }
      this.setState({ project: newProject })
    })
    en(this.getCosts)
    en(err => console.error(err))
  }
}

const Cost = (cost) => {
  const { project_id } = cost
  const { costs } = cost
  const cost_id = costs[cost.id]
  const cost = costs[cost_id]
  const cost_index = costs.indexOf(cost)
  const newCosts = [...costs]
  newCosts[cost_index] = cost
  newCosts[cost_index].isDeleted = true
  newCosts[cost_index].deletedAt = Date.now()
  newCosts[cost_index].cost_id = cost_id
  const newProject = { ...cost.project, costs: newCosts }
  return (
    <a href={cost.project_id}>
      <img alt="Delete icon" data-bbox="100 100 120 120" />
      <span>Delete</span>
    </a>
  )
}
```

October 21th, 2016

Dyn DDoS Attack

Twitter, Netflix, and Reddit



July 19th, 2024

Microsoft's fall

**IT Outage for 8.5 mill devices
(Including Sonia)**



December 9th, 2020

**European Medicines
Agency**

INTRODUCTION



- Network intrusions pose **significant threats**
- Automated detections using ML offers **scalable solutions**

Traditionally, detection methods are **rule based and reactive**

Our Goal: Prepare and Develop a network intrusion detection system using the UNSW-NB15 dataset, so as to prepare for the attacks

DATASET FEATURES



UNSW
THE UNIVERSITY OF NEW SOUTH WALES

UNSW-NB15 DATASET

AVAILABLE IN KAGGLE

- ~250,000 network traffic records (~180,000 in our training set)
- Consisting of normal traffic and 9 different types of attacks
- 45 features include basic, content-based, time-based, and traffic-based features

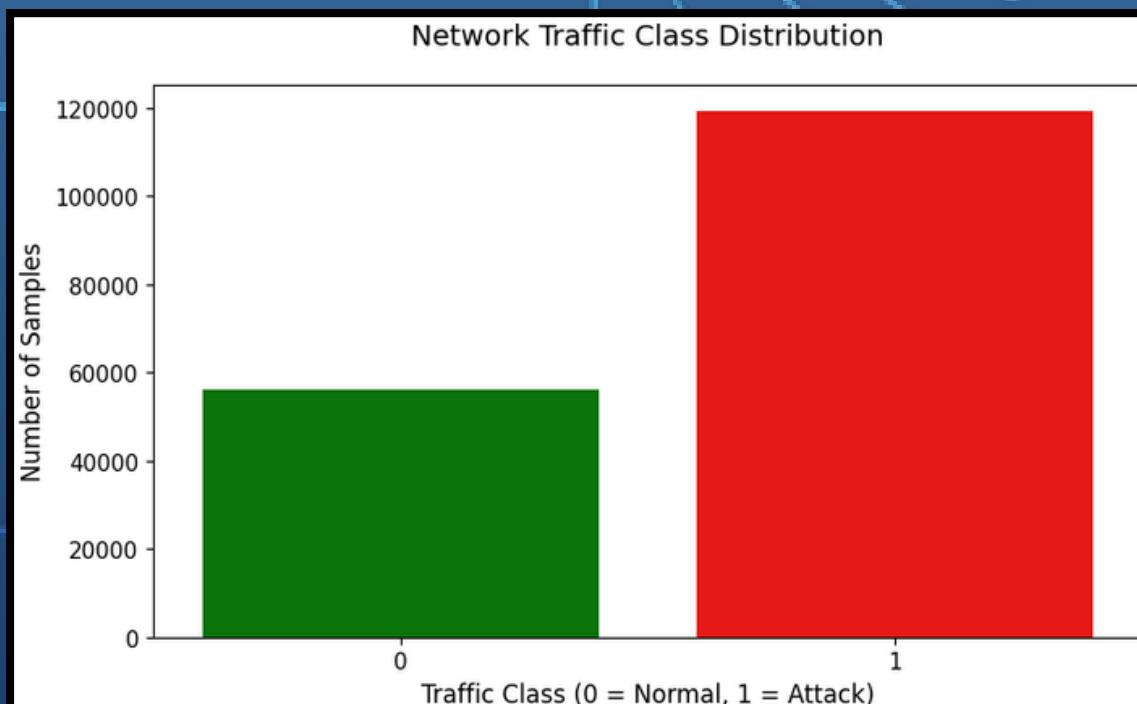
ACSC Australian
Cyber Security
Centre

DEVELOPED BY THE ACSC

LABELED NORMAL AND MALICIOUS TRAFFIC PATTERNS

9 attack categories

Fuzzers, Analysis,
Backdoors, DoS,
Exploits, Generic,
Reconnaissance,
Shellcode, and Worms



RE



QUESTIONS

-  Can **PCA** improve model performance?
-  How well do **K-Means** and **KNN** distinguish malicious traffic?
-  Can **collaborative filtering** predict threats based on historical patterns?
-  What are the **trade-offs** of each analyzed technique?



PREPROCESSING STEP

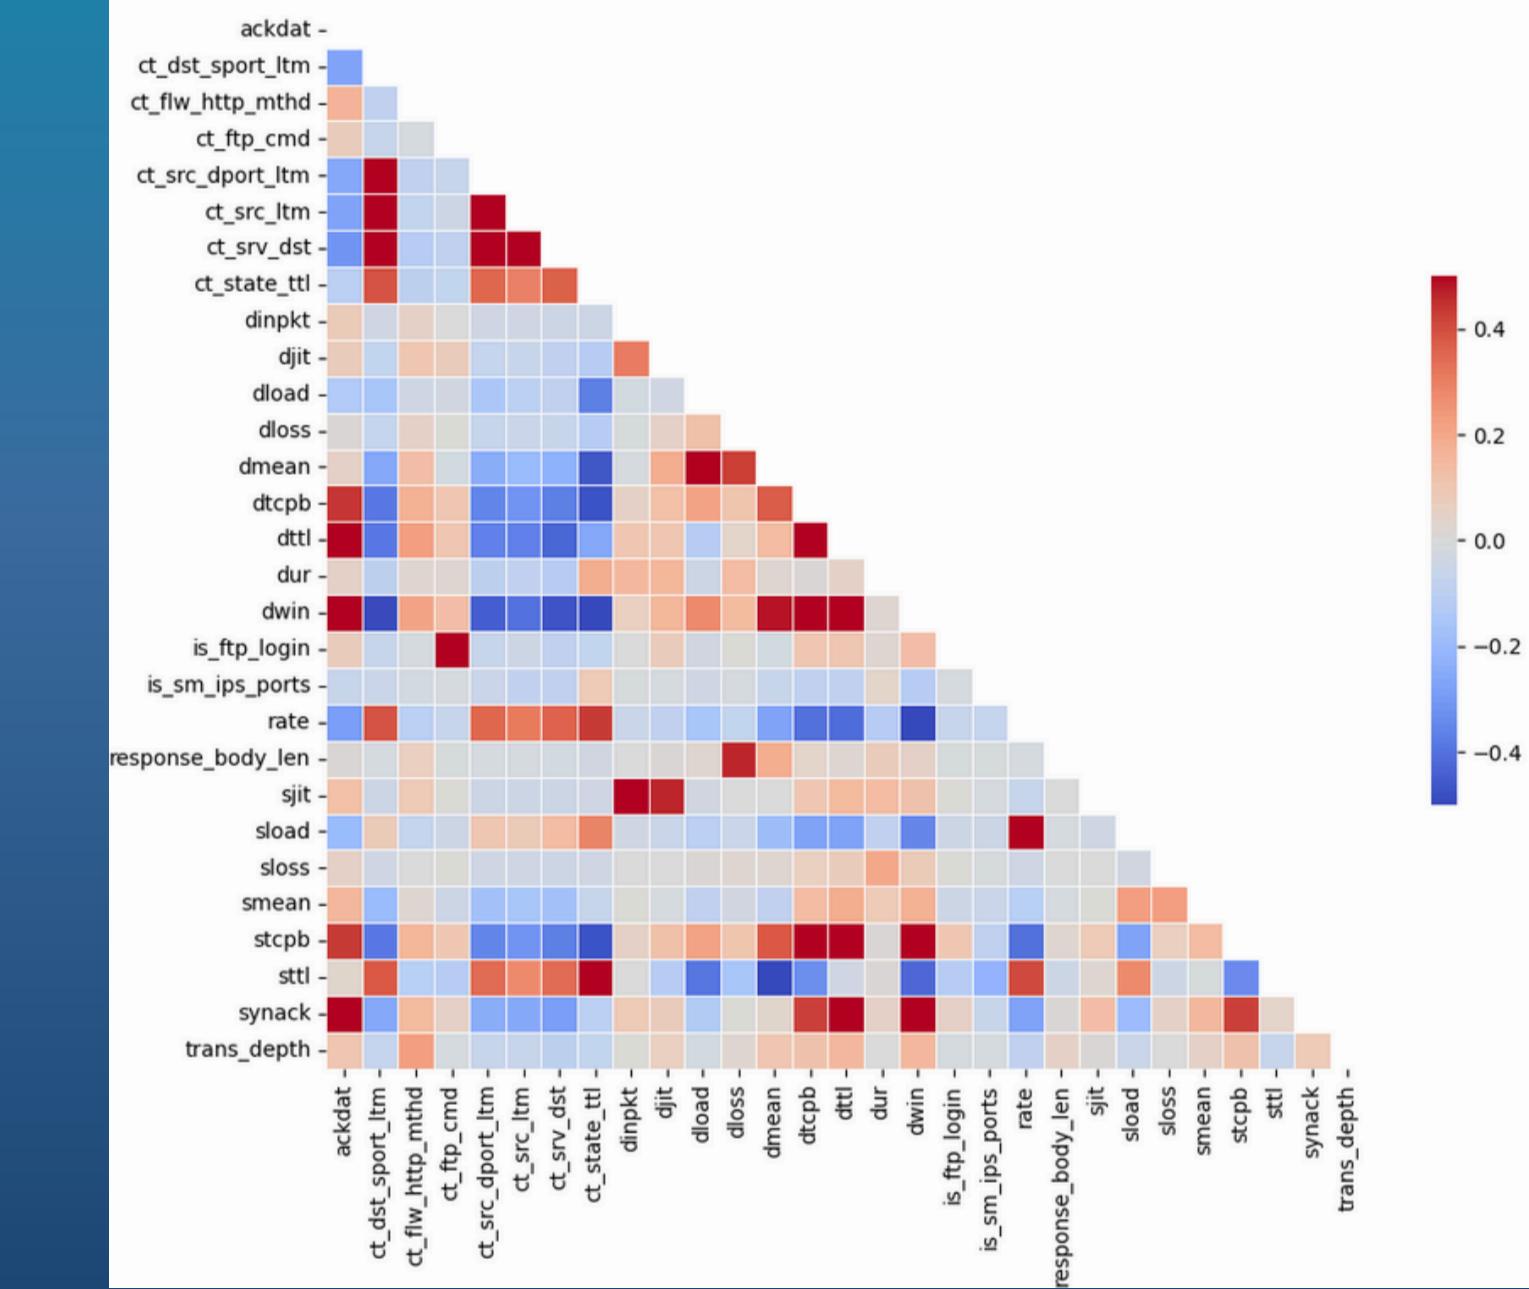
CLEANING

- Remove irrelevant features like attack category, label number, and id number (no insight to training)
- Drop variables that have high correlation with each other
- 29 features remain to be trained and tested

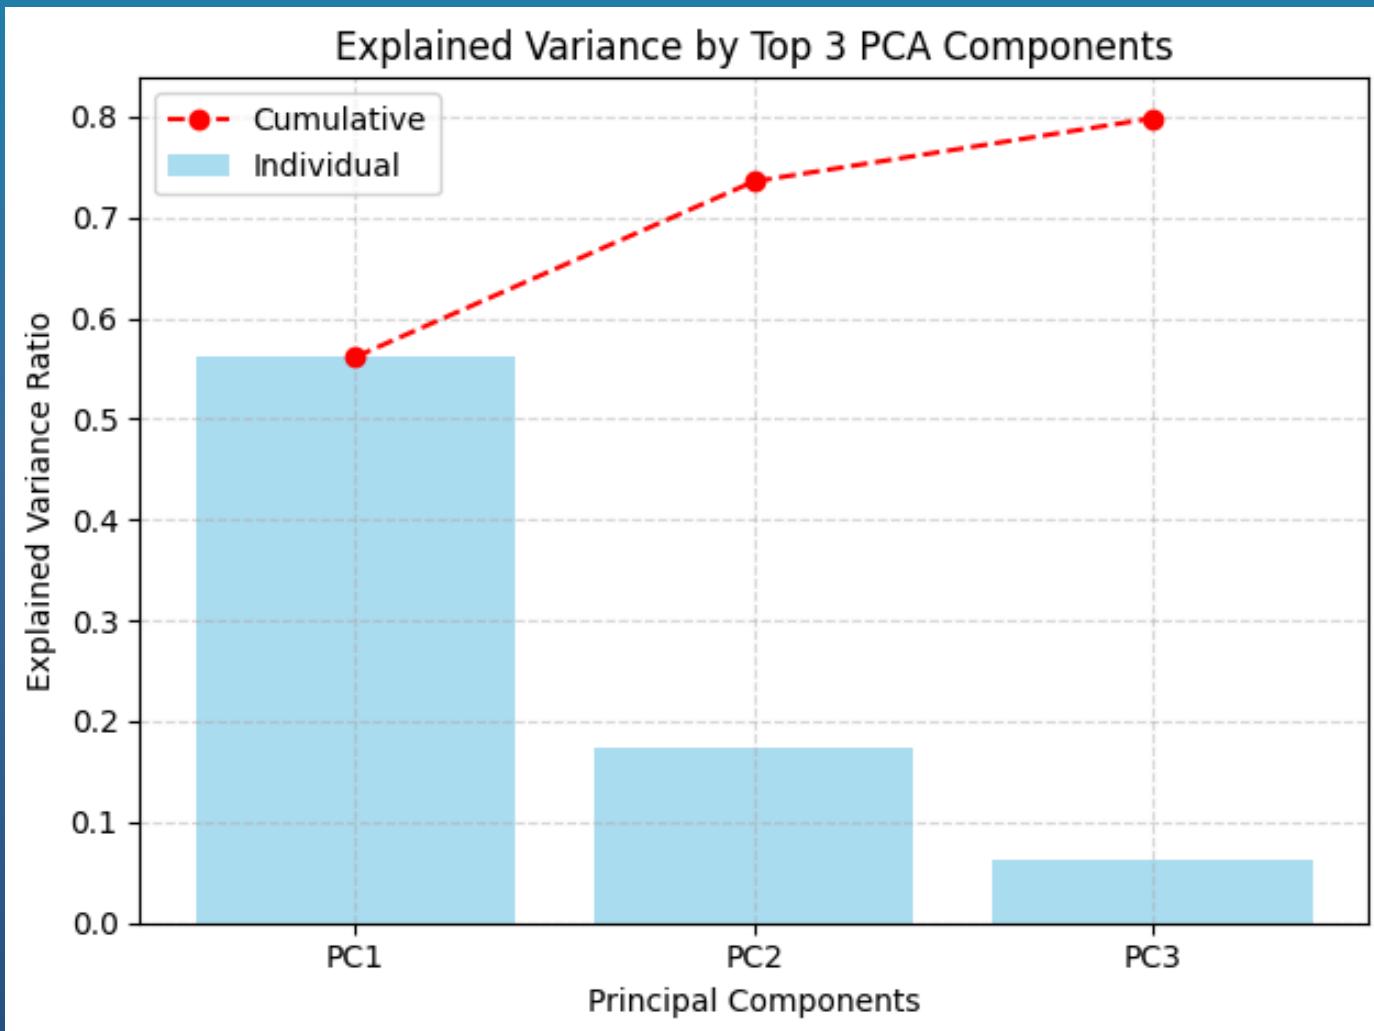
SCALING

Normalize the data to perform PCA

- StandardScaler. Centers data ($\text{mean}=0$, $\text{std}=1$).
- **MinMaxScaler**. Scales to $[0, 1]$ range



PRINCIPAL COMPONENT ANALYSIS

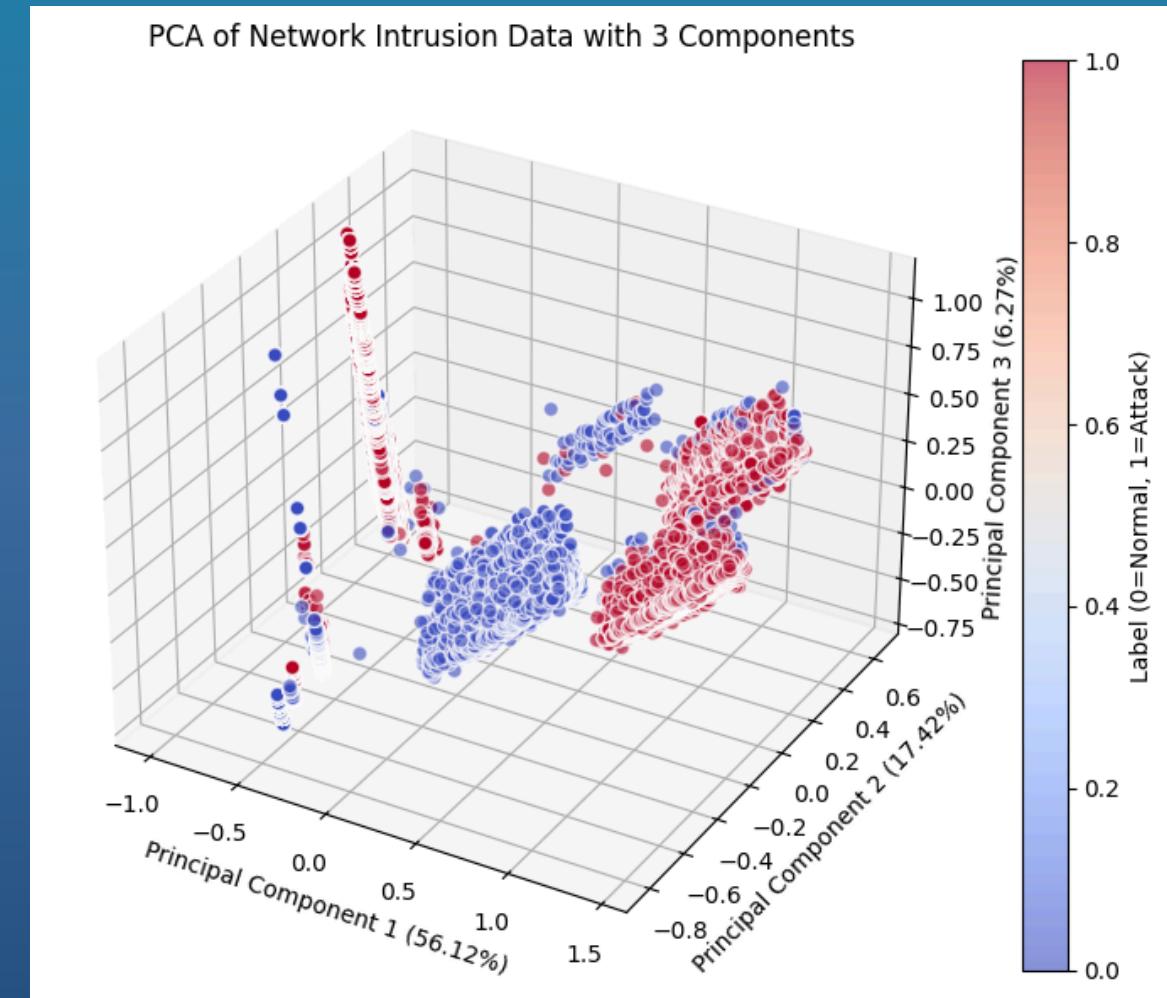


USED:

min-max scaled matrix
due to skewness in the data

CHOICE:

3 feature components
over initial 2

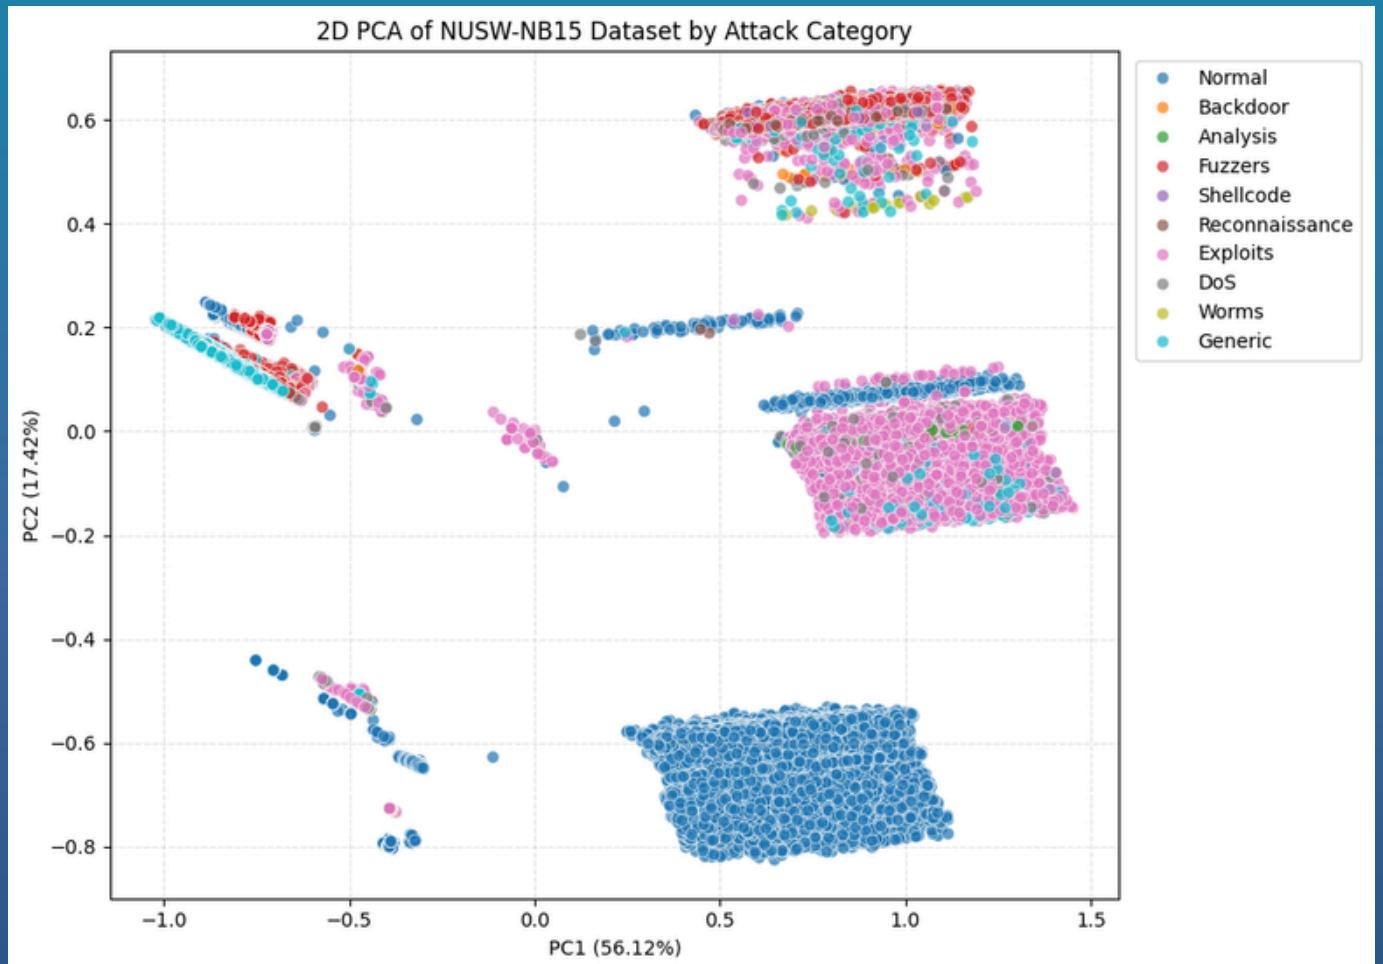


Numerical columns reduced: from 39 to 29

PCAs consist of most deterministic variables

Categorical variables (strings): *proto*, *service* and *state*

SUPERVISED LEARNING ALGORITHM



Data in PCA shows clear groups between attacks and not

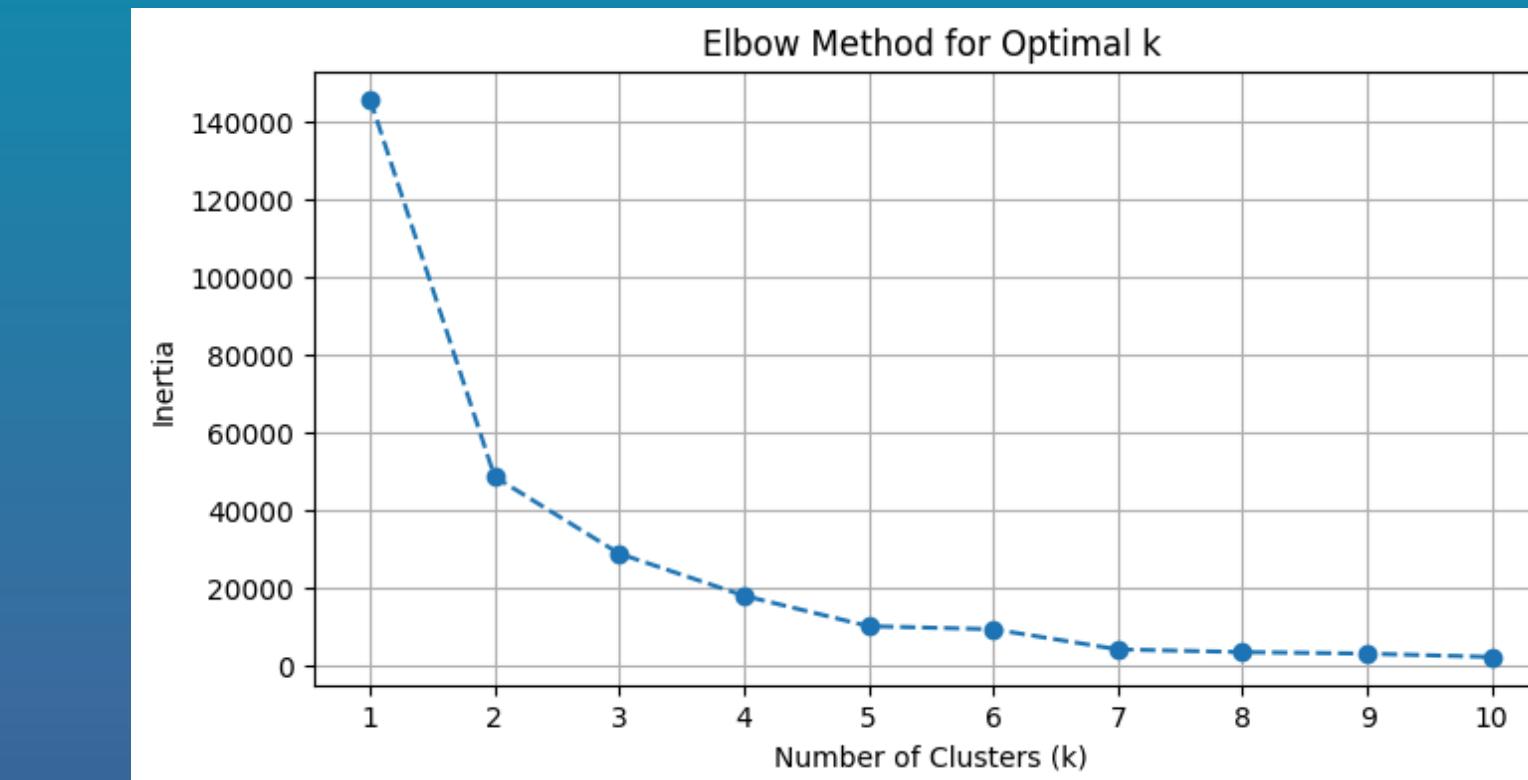
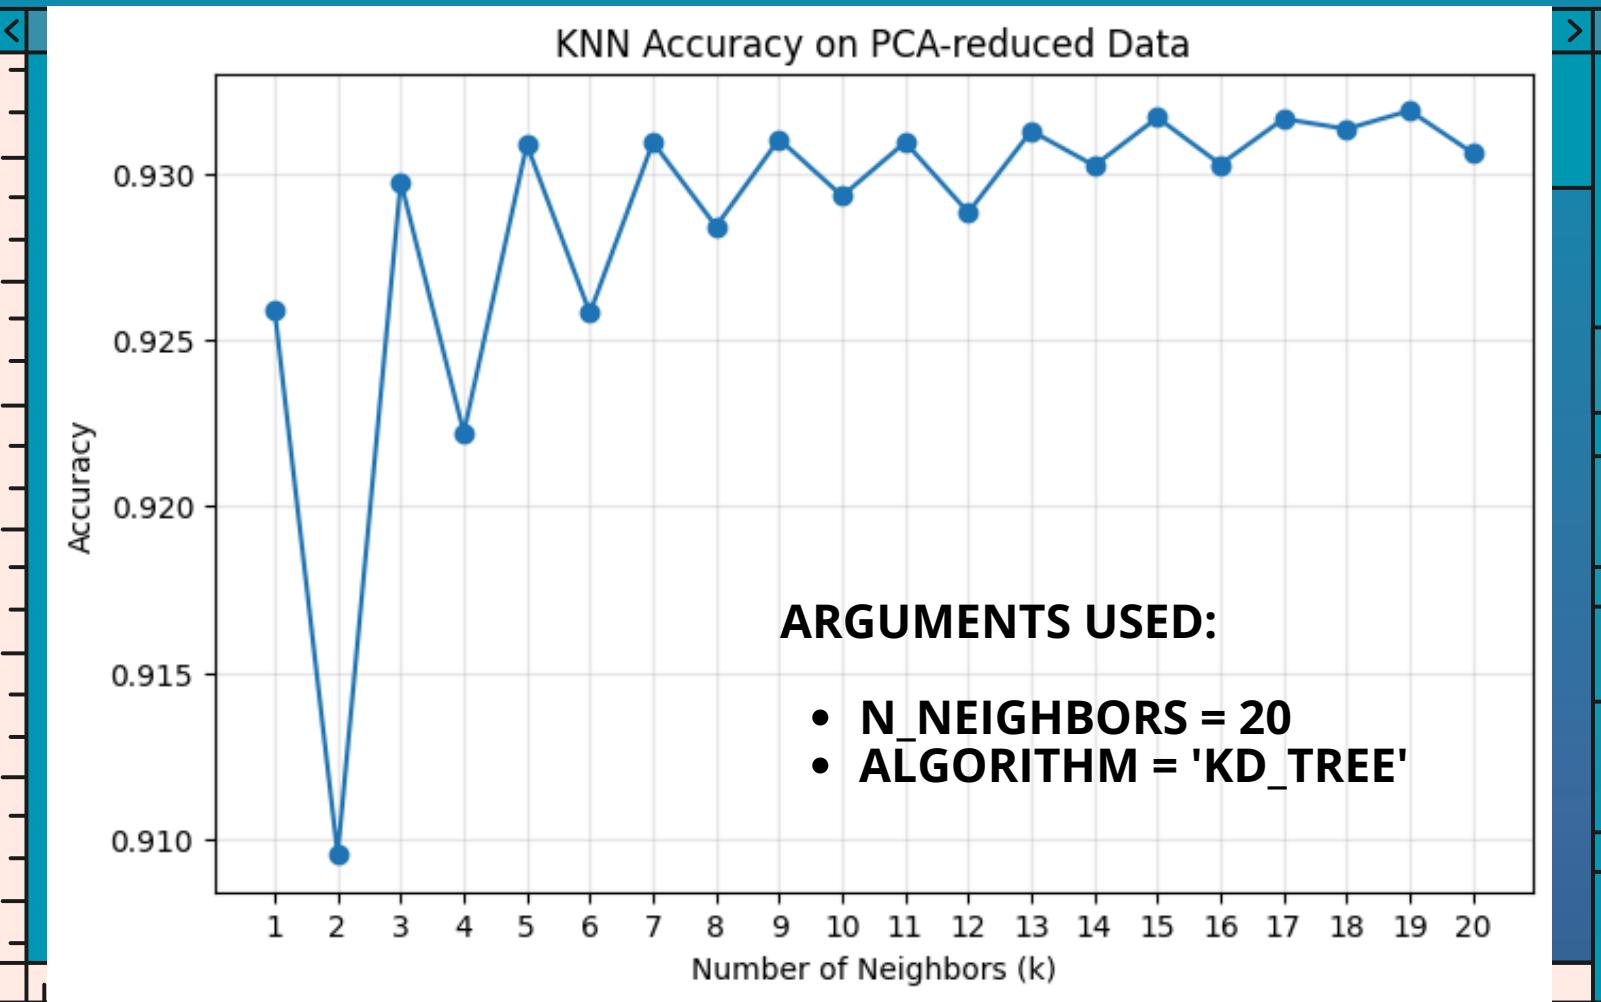
K-MEANS CLUSTERING



K-means does not improve on PCA separation, but rather just separates normal traffic from attacks

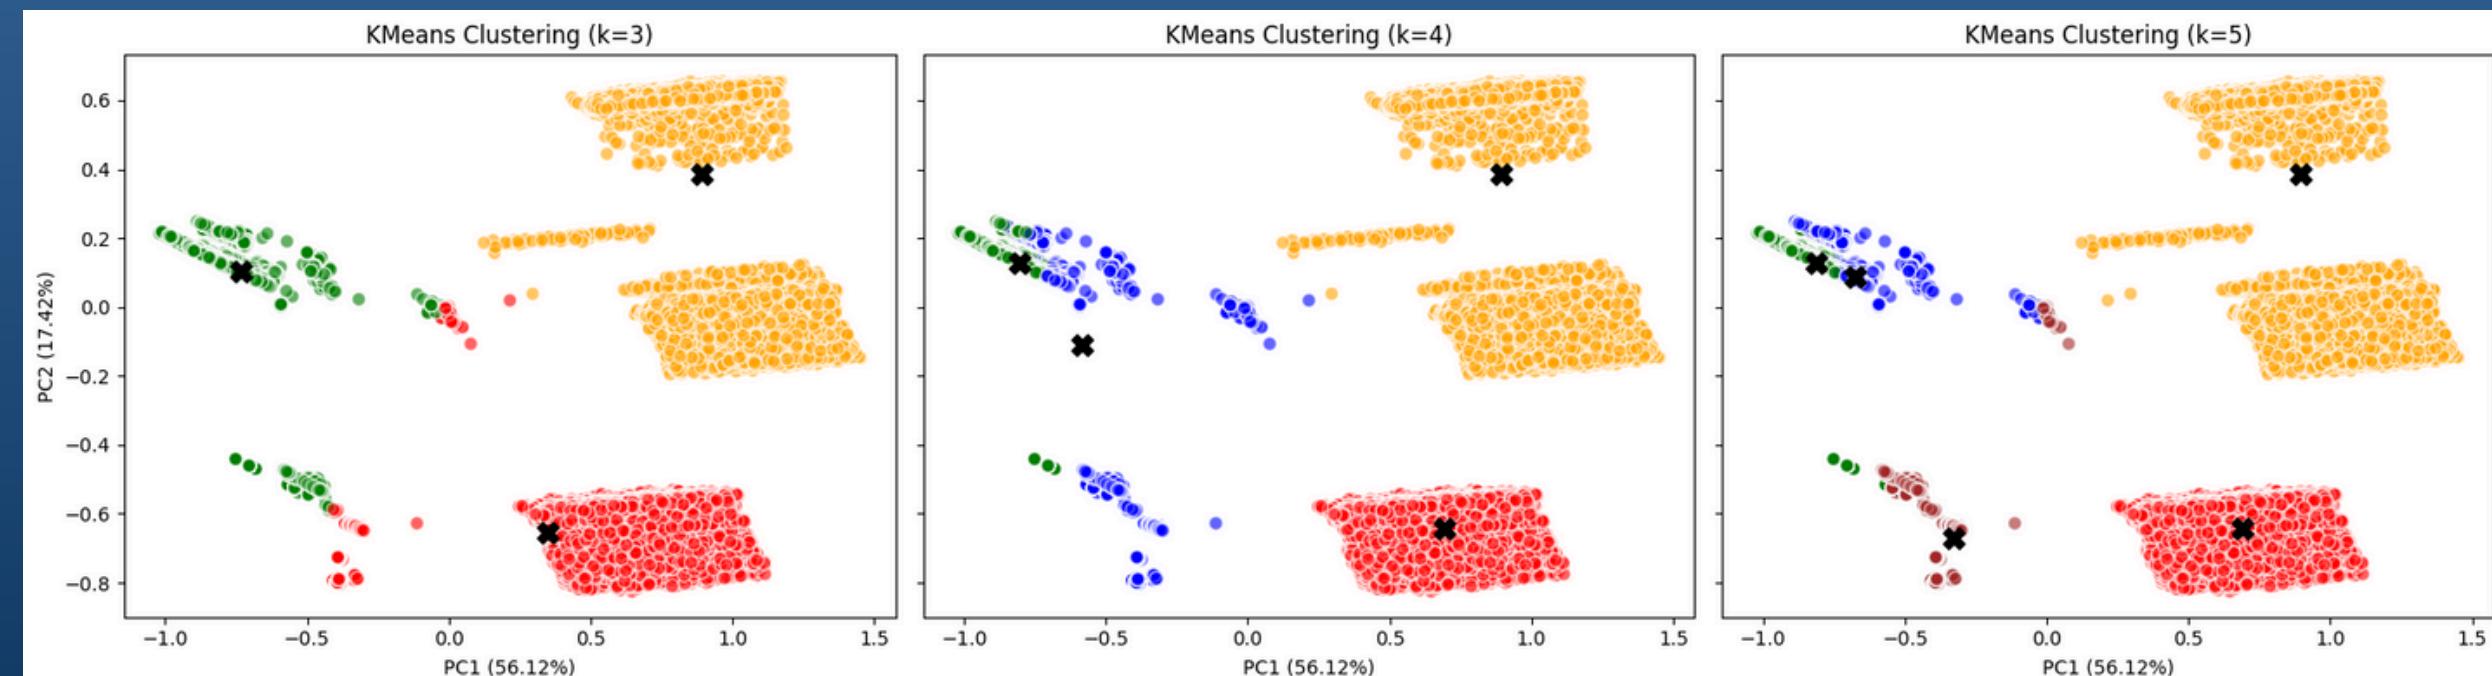
CLASSIFYING

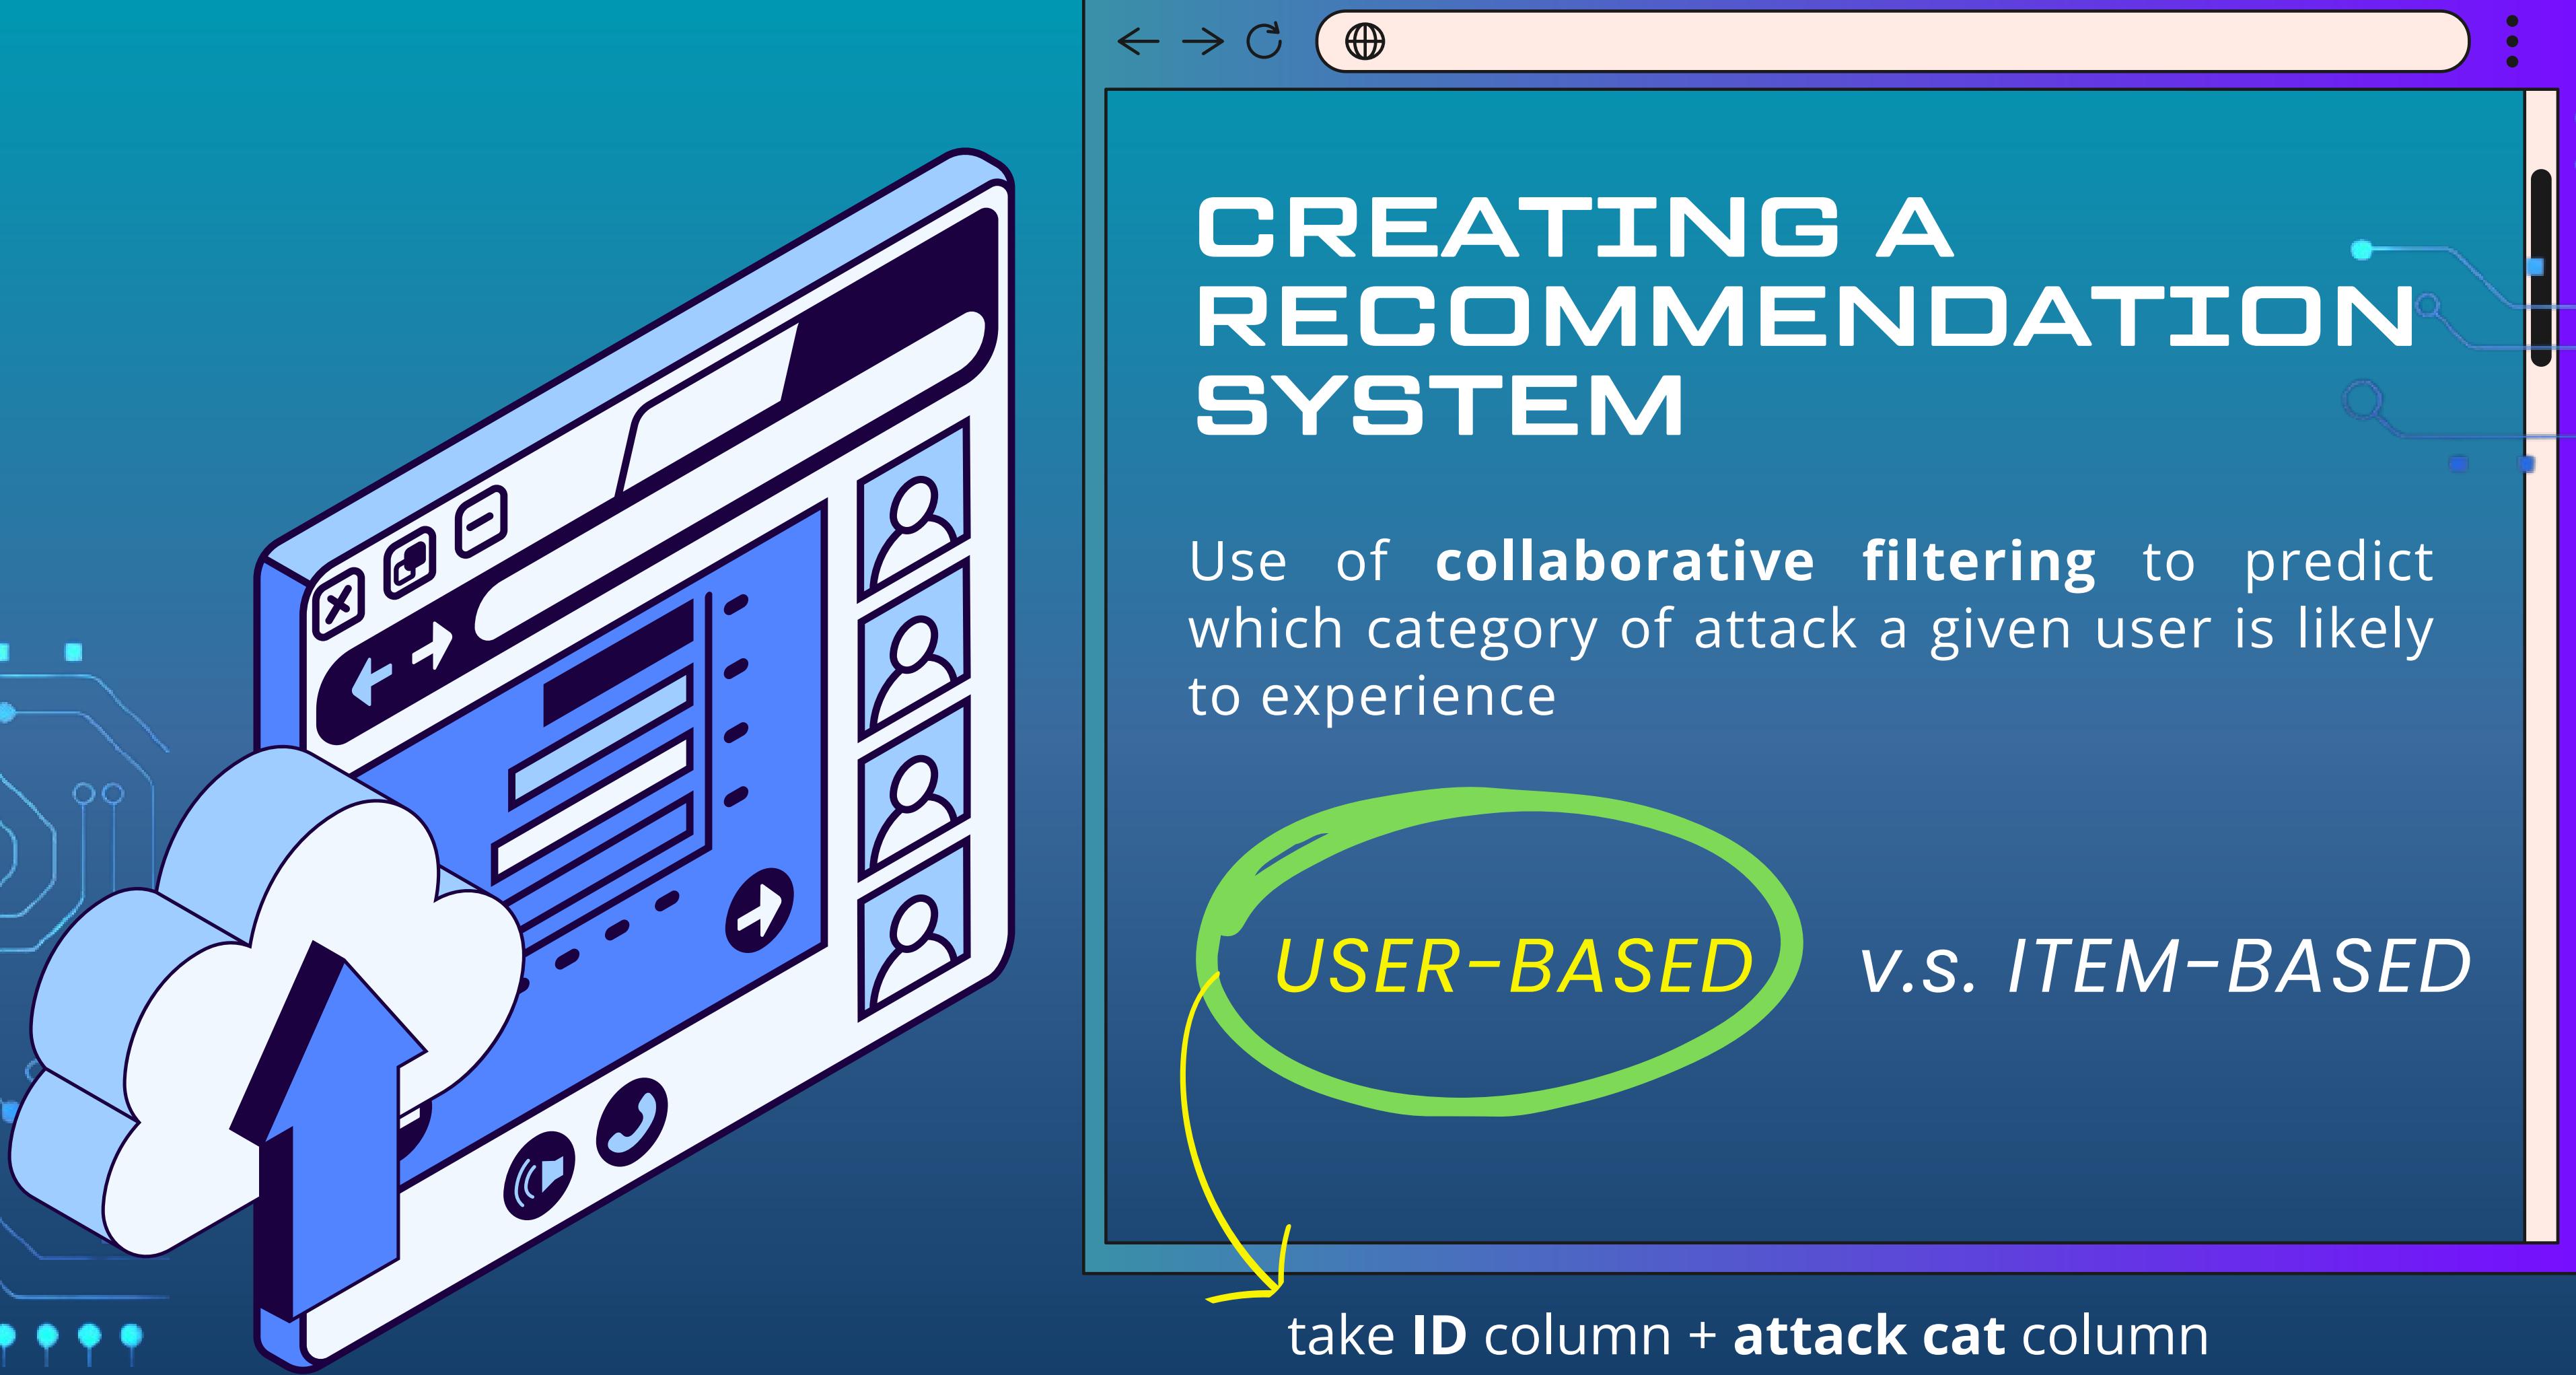
K-NEAREST NEIGHBORS



ELBOW METHOD INDICATES $k \approx [3,5]$,
LIKELY OPTIMAL CLUSTER RANGE

BELOW SHOWS RANGE OF KMEANS ON PCA-REDUCED DATA



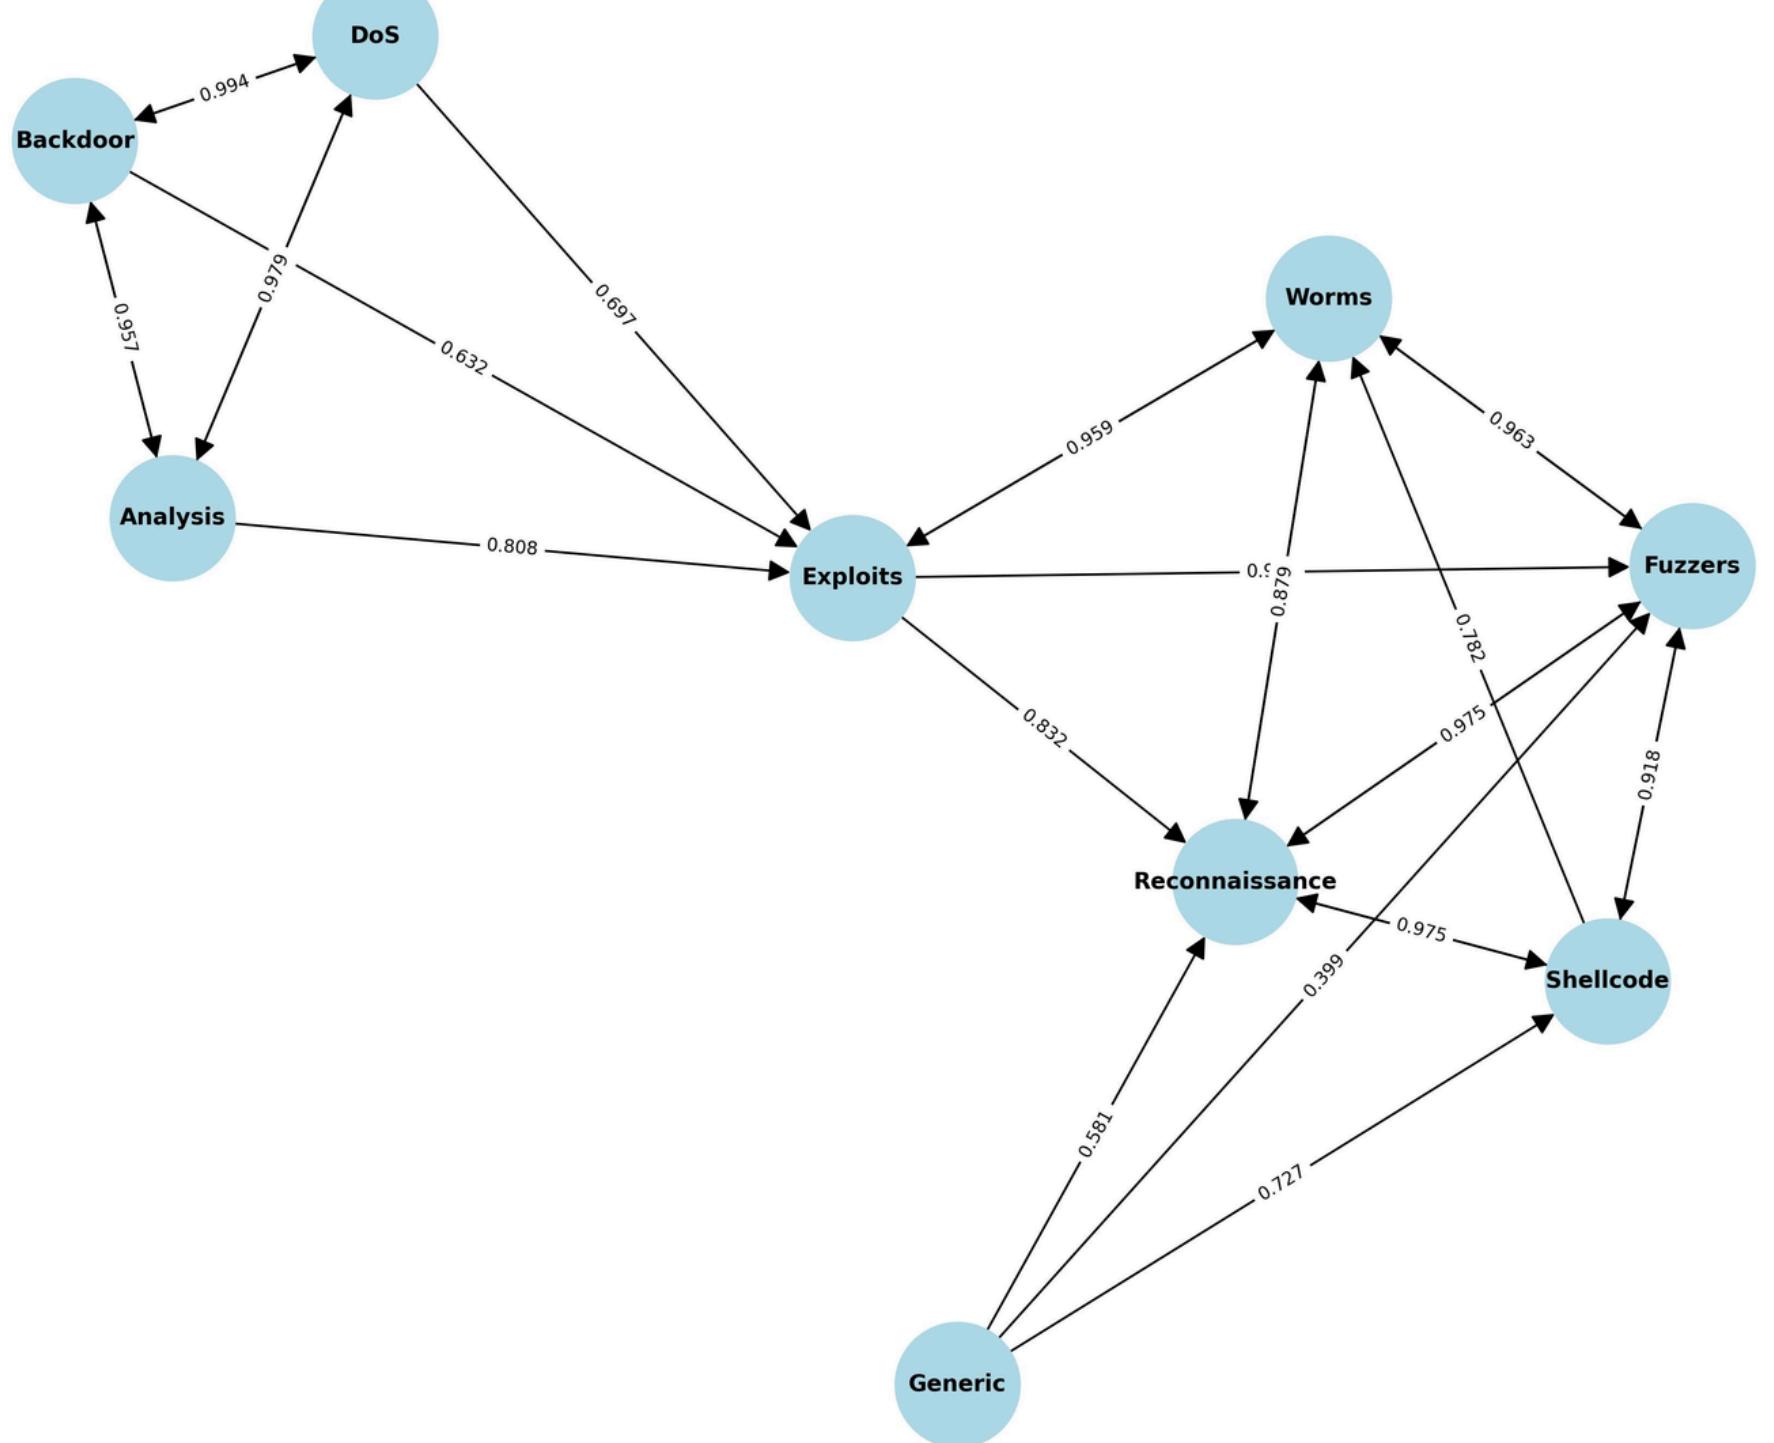


category distribution

| | count | percentage |
|----------------|---------------|------------|
| attack_cat | | |
| Normal | 56000 | 31.94 |
| Generic | 40000 | 22.81 |
| Exploits | 33393 | 19.04 |
| Fuzzers | 18184 | 10.37 |
| DoS | 12264 | 6.99 |
| Reconnaissance | 10491 | 5.98 |
| Analysis | 2000 | 1.14 |
| Backdoor | 1746 | 1.00 |
| Shellcode | 1133 | 0.65 |
| Worms | 130 | 0.07 |
| TOTAL: | 175341 | |

9
ATTACK TYPES

attack similarity

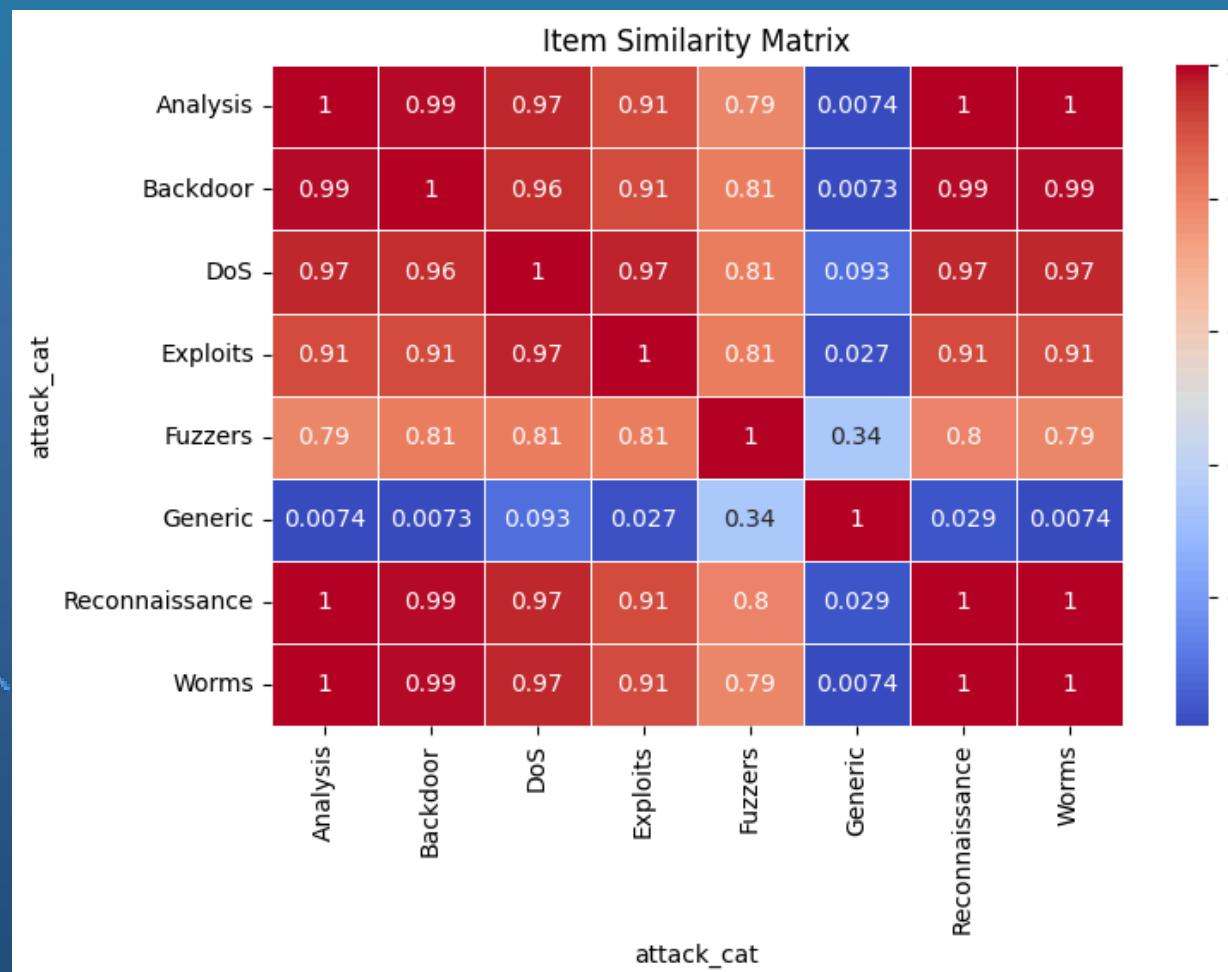
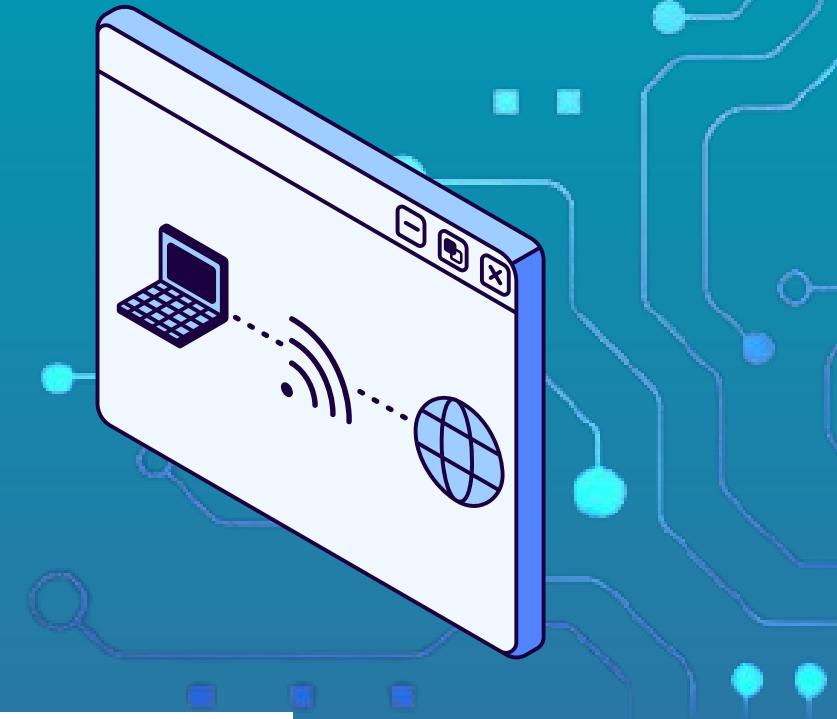


label column

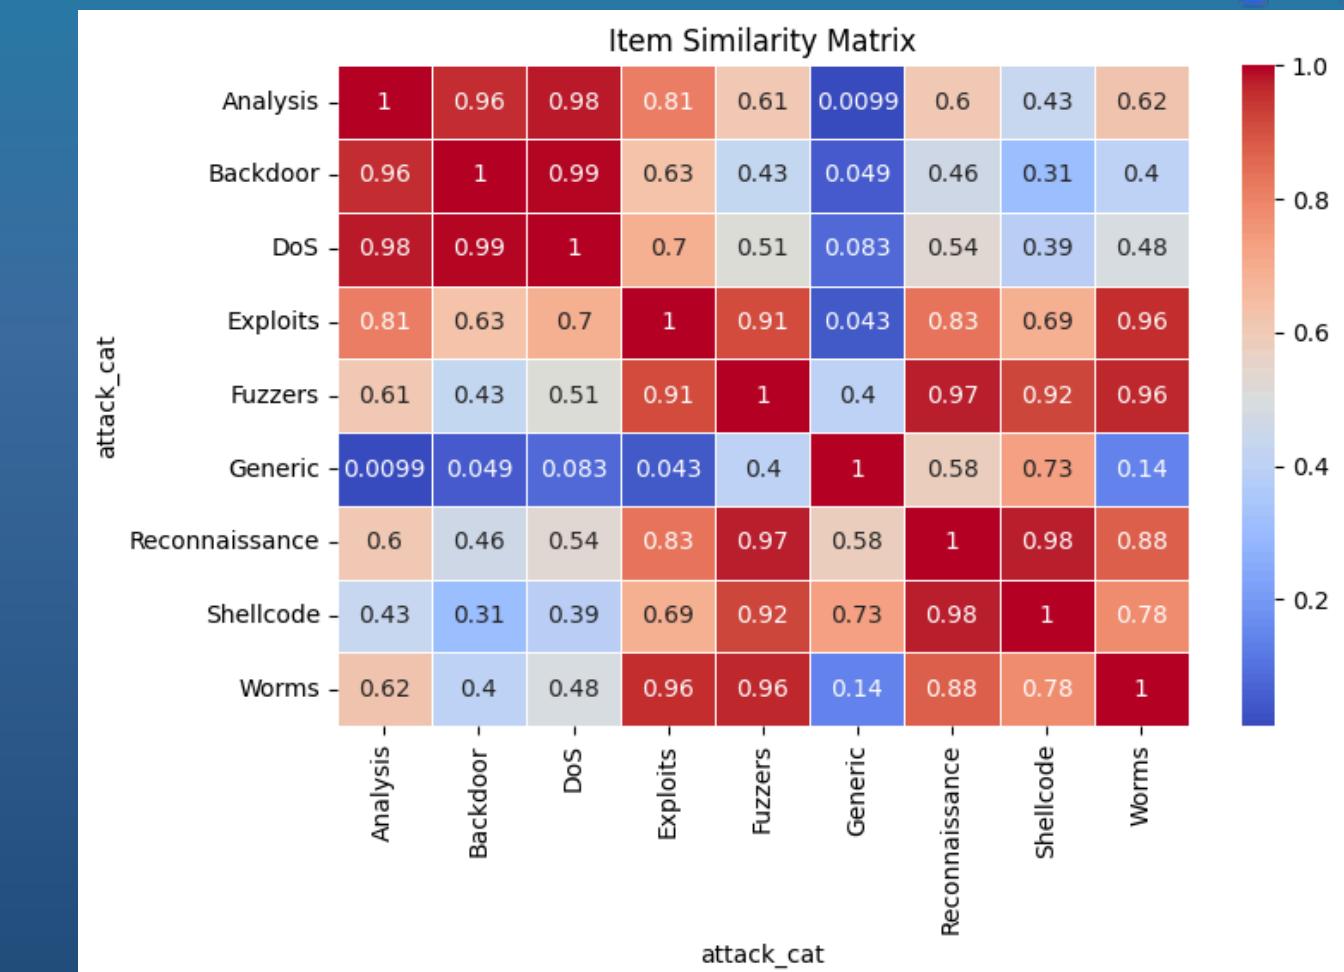
NOT ATTACKED (0) OR ATTACKED (1)

'Similar users experience similar attack patterns'

Cosine similarity makes difference in the prediction accuracy in **multi-class classification** (attack category).



based on service usage patterns



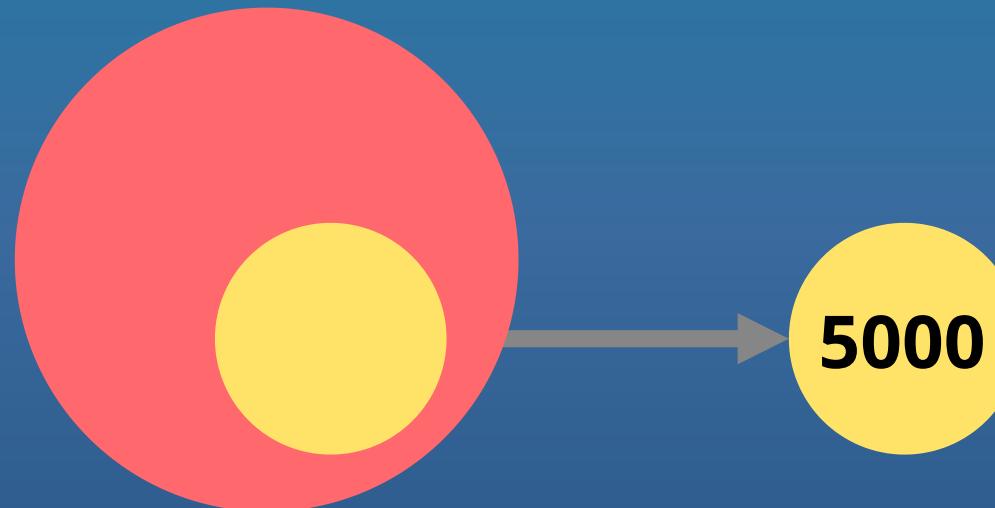
based on transaction protocols

STRATIFIED SAMPLING

LABELS: NORMAL (=0) AND MALICIOUS TRAFFIC (=1)

TWO STRATAS BASED ON SHARED CHARACTERISTICS

175341 X 175341
(matrix size)



RANDOM SAMPLE
OF 5000 USERS



NORMAL TRAFFIC

Consider the
worst scenario

'UNKNOWN'
USERS

Label 0

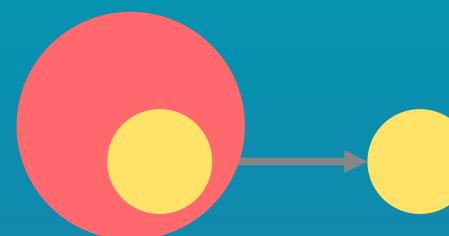


MALICIOUS TRAFFIC

ATTACKED
KNOWN USERS

Label 1

Users with no attack (0=Normal) are the **targets** for the later prediction



SAMPLE SIMILARITY MATRIX

Sampled user similarity matrix shape: (5000, 5000)

| id | 15483 | 133350 | 80486 | 29973 | 18340 | 170501 | 165831 | 55216 | 37757 | 22161 | ... | 93512 | 154465 | 146874 | 73134 | 97379 | 57479 | 145823 | 9277 | 1091 | 6986 |
|--------|-------|--------|-------|-------|-------|--------|--------|-------|-------|-------|-----|-------|--------|--------|-------|-------|-------|--------|------|------|------|
| id | | | | | | | | | | | | | | | | | | | | | |
| 15483 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 133350 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 1.0 | 1.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 1.0 | 1.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 |
| 80486 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 1.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 29973 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 18340 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 57479 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 145823 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 1.0 | 1.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 1.0 | 1.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 |
| 9277 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 1091 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 6986 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |

5000 rows × 5000 columns



NEAREST NEIGHBORS FOR PREDICTING CYBER THREADS

1. Loop through target users and find their nearest neighbors
2. Compare them to known users to infer likely attack types
3. Recommend the most common attack type from those neighbors



closest neighbors RESULTS

Random selected target ID

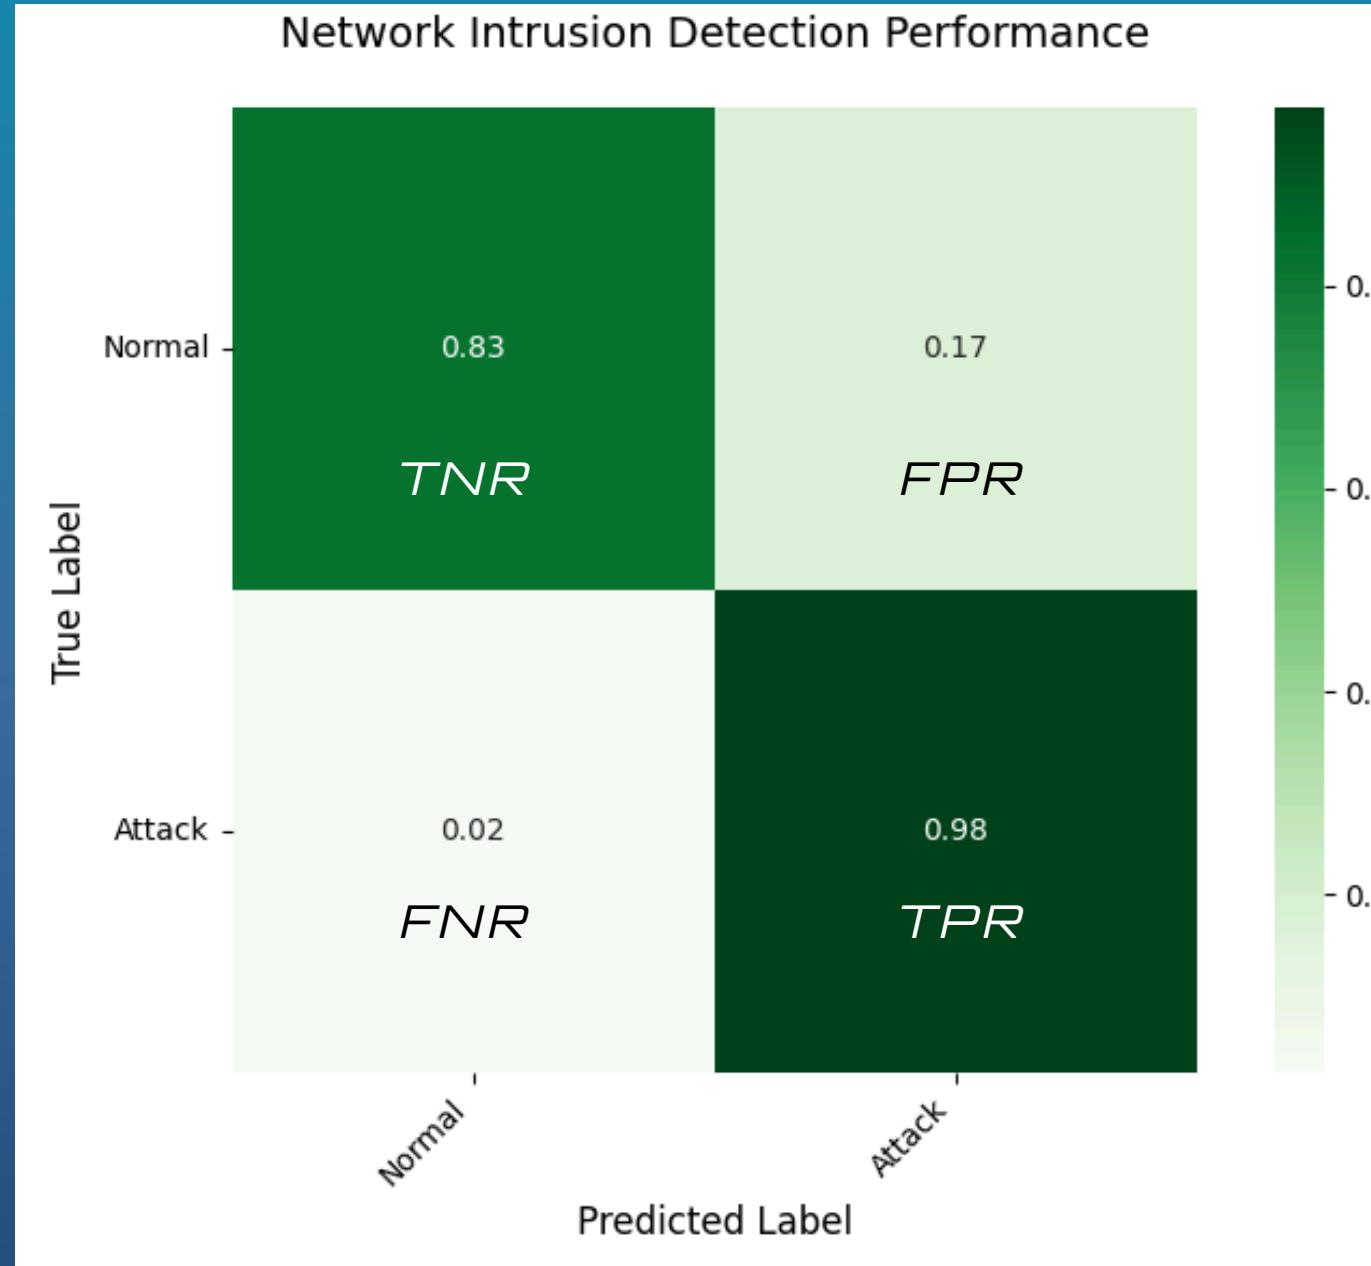
```
target_id =  
random.choice(similarity_df.i  
ndex)
```

| Attack category counts: | 0 |
|-------------------------|---|
| attack_cat | 0 |
| Analysis | 0 |
| Backdoor | 0 |
| DoS | 0 |
| Exploits | 0 |
| Fuzzers | 0 |
| Generic | 0 |
| Reconnaissance | 0 |
| Shellcode | 5 |
| Worms | 0 |

| Target ID | Top Neighbors | Attack Category Counts | Predicted Attack Category |
|-----------|--------------------------------------|--------------------------|---------------------------|
| 0 118838 | [83977, 52190, 91316, 82963, 145861] | [0, 0, 0, 0, 0, 0, 5, 0] | Shellcode |

Prediction for that ID

RESULTS. PERFORMANCE METRICS.



CLASSIFICATION REPORT AFTER PREDICTIONS

| | Normal | Attack | accuracy | macro avg | weighted avg |
|-----------|--------------|--------------|----------|--------------|--------------|
| precision | 0.943791 | 0.925461 | 0.930612 | 0.934626 | 0.931305 |
| recall | 0.831922 | 0.976808 | 0.930612 | 0.904365 | 0.930612 |
| f1-score | 0.884333 | 0.950441 | 0.930612 | 0.917387 | 0.929363 |
| support | 16772.000000 | 35831.000000 | 0.930612 | 52603.000000 | 52603.000000 |

↙ # actual occurrences

TNR: ~83% of the 16,772 normal users were correctly identified

High TPR: ~98% of actual attacks were correctly detected

harmonic mean The model balances precision/recall better for attacks than normal traffic ($F1\ score = 2pr / (p+r)$)

CONFUSION MATRIX

NORMAL AS - CLASS, ATTACK AS + CLASS

FP (FALSE ALARMS) > FN (MISSED ATTACKS)

CONCLUSION. PROS AND CONS.

pca>

- Reduces noise, speeds up training (~80% variance in 2-3 PCs)
- Loses interpretability

<k-means

- Groups similar traffic (unsupervised)
- Needs manual labeling, sensitive to noise

knn>

- 93% accuracy, easy to implement
- Slow on large data, imbalance-sensitive

<CF

- Predicts co-occurring attacks (e.g., Exploits after DoS)
- Needs rich history, as it struggles with new threats

Best for UNSW-NB15: PCA + KNN (balance of speed & accuracy)



THANK YOU FOR YOUR ATTENTION